

REHAN AHMAD

BS Artificial Intelligence — 6th Semester, University of Wah

0342-4983463 | m.rehanahmad25@email.com

[rehanahmad.dev](mailto:m.rehanahmad25@email.com) | github.com/rehan-ahmad25 | linkedin.com/in/rehan-ahmad-ai

CAREER OBJECTIVE

Motivated BS Artificial Intelligence student (6th Semester) with hands-on experience shipping end-to-end AI applications spanning LLMs (built a GPT-style Transformer from scratch), agentic voice AI, computer vision, and explainable deep learning. Comfortable across the full stack — from model training to FastAPI/Flask deployment. Seeking an AI/ML internship to contribute to production-grade AI systems while deepening research and engineering skills.

EXPERIENCE

AI/ML Intern — Redfort Technologies

July 2026 – Present

- Working on Transformer-based architectures, contributing to the development and experimentation of deep learning models for real-world AI applications.
- Collaborating with the engineering team on model implementation, evaluation, and optimization within a production AI workflow.

TECHNICAL SKILLS

- Languages:** Python, C++, Java, HTML, CSS, JavaScript
- AI / ML:** Machine Learning, Deep Learning, Neural Networks (ANN, CNN, RNN), Computer Vision
- Libraries & Frameworks:** NumPy, Pandas, Scikit-learn, TensorFlow, Keras, PyTorch, Flask, FastAPI
- AI Tooling:** YOLOv8, MediaPipe, OpenAI Whisper, FAISS, LangChain/Ollama for RAG pipelines
- Tools:** Jupyter Notebook, Google Colab, VS Code, Git/GitHub

FEATURED PROJECTS (GITHUB.COM/REHAN-AHMAD25)

AstraGPT — GPT-Style LLM Built From Scratch | *PyTorch, Transformers, Self-Attention*

<https://github.com/rehan-ahmad25/AstraGPT>

- Implemented a decoder-only Transformer language model from first principles — custom tokenizer, multi-head self-attention, positional embeddings, and training loop — with no pre-trained weights or external APIs.
- Built the full LLM pipeline end-to-end (data preprocessing to autoregressive generation), demonstrating deep understanding of the architecture behind modern GPT-style models.

Astra Voice Receptionist — AI Voice Agent for Healthcare | *Vapi, FastAPI, SQLAlchemy, SQLite*

<https://github.com/rehan-ahmad25/Astra-voice-receptionist>

- Built a production-style voice AI receptionist for dental clinics that handles appointment booking, rescheduling, cancellations, and patient inquiries through natural conversation.
- Designed a FastAPI backend with SQLAlchemy-modeled clinic data, integrated with Vapi's voice pipeline for real-time conversational automation.

Explainable Brain Tumor Classification | *PyTorch, EfficientNet-B0, Grad-CAM, SHAP, FastAPI*

<https://github.com/rehan-ahmad25/explainable-brain-tumor-classification>

- Built a multi-class MRI brain tumor classifier using EfficientNet-B0, paired with Grad-CAM and SHAP to make model predictions visually and quantitatively explainable.
- Served predictions and explanations via a FastAPI backend, prioritizing model transparency for a high-stakes medical use case.

AI-Powered Crime Analyzer | *Python, YOLOv8, OpenAI Whisper, Flask*

<https://github.com/rehan-ahmad25/AI-Powered-Crime-Analyzer>

- Built a multi-modal forensic analysis platform combining a custom fine-tuned YOLOv8 model for weapon detection with OpenAI Whisper for audio transcription.

- Designed a secure pipeline to analyze video and audio evidence together, surfacing weapon detections alongside transcribed spoken content for investigators.

LectureFlix AI — AI Learning Platform | *Flask, Whisper, FAISS, Ollama, LLM APIs*

<https://github.com/rehan-ahmad25/LectureFlix-AI>

- Built a full-stack platform that converts lecture videos, YouTube content, and PDFs into searchable transcripts, summaries, flashcards, and quizzes.
- Integrated Whisper for transcription and FAISS-based semantic search with local LLMs (Ollama) for intelligent, context-aware Q&A.

Eye Disease Detection | *PyTorch, EfficientNet-B0, Flask*

<https://github.com/rehan-ahmad25/eye-disease-detection-ai>

- Developed a deep learning system that detects Cataract, Glaucoma, and Diabetic Retinopathy from retinal fundus images, extending from binary to multiclass classification.
- Applied data augmentation and preprocessing to correct class imbalance, achieving 90%+ classification accuracy, and deployed as a Flask web app.

EDUCATION

- **BS Artificial Intelligence — University of Wah:** 2023 – Present | CGPA: 3.9 / 4.0 | Expected Graduation: 2027
- **FSc Pre-Engineering — FBISE:** 2021 – 2023 | 936/1100
- **SSC (Science) — FBISE:** 2019 – 2021 | 1016/1100

RESEARCH INTERESTS

Large Language Models (LLMs) and AI-powered reasoning systems • Retrieval-Augmented Generation (RAG) and tool-augmented AI • Computer Vision and multi-modal AI systems • AI for public safety, healthcare, and financial anomaly detection • Generative AI and intelligent agent systems

LANGUAGES

Urdu — Native | English — Proficient